

# Lecture 10: Topic Modeling, Word2Vec

LING-351 Language Technology and LLMs

---

Instructor: Hakyung Sung

September 25, 2025

# Table of contents

---

1. Word distributions

2. Word vectors

# Review

---

- Tokenization
- Lemmatization
- Frequency calculation
- Concordance
- (Collocation)

# Word distributions

---

Examining word distributions is the first and most important step in corpus/text analysis.

# Zipf's power law (1932)

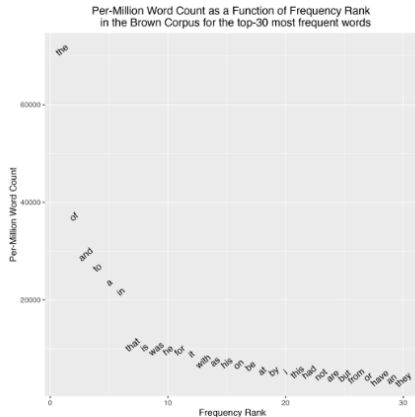


Figure 4.3: Per-million-word frequency of words in the Brown Corpus as a function of their frequency rank (ordered from left to right as the first most frequent word, the second most frequent, and so on).

# Zipf's power law (1932)

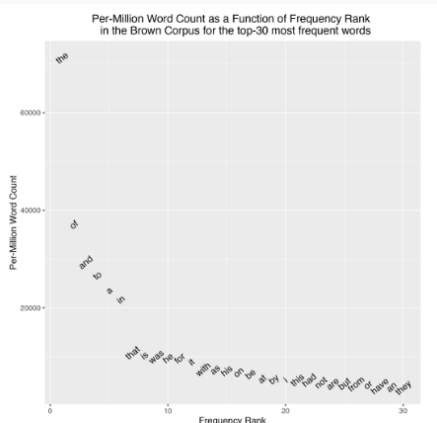


Figure 4.3: Per-million-word frequency of words in the Brown Corpus as a function of their frequency rank (ordered from left to right as the first most frequent word, the second most frequent, and so on).

**Implication:** Few words are very frequent; many are rare  $\Rightarrow$  long tail.



# Zipf's power law (1932)

- type vs. token

# Zipf's power law (1932)

- type vs. token
- frequency  $\propto 1/\text{rank}$ .

# Zipf's power law (1932)

- type vs. token
- frequency  $\propto 1/\text{rank}$ .
- e.g., Brown corpus: *the*  $\approx 6\%$  tokens; *of*  $\approx 3\%$ ; *and*  $\approx 2.6\%$ .

# Zipf's brevity law




Figure 4.4: Histogram of the length (number of characters) of all word types in the Brown Corpus.

- More frequent words tend to be **shorter** (characters/syllables).

- More frequent words tend to be **shorter** (characters/syllables).
- Efficiency pressure: frequent items economize articulatory/processing effort.

- More frequent words tend to be **shorter** (characters/syllables).
- Efficiency pressure: frequent items economize articulatory/processing effort.
- Most frequent Brown words: monosyllabic,  $\leq 3$  letters (*the, of, and, a, in, to, is, was, I, for*).

# Heaps' law

As you read more tokens in a corpus, you keep seeing new word **types**, but the **rate** of new words **slows down**.



# Heaps' law

As you read more tokens in a corpus, you keep seeing new word **types**, but the **rate** of new words **slows down**.

## Example

- After 1,000 tokens: ~700 unique words

Why it matters?

# Heaps' law

As you read more tokens in a corpus, you keep seeing new word **types**, but the **rate** of new words **slows down**.

## Example

- After 1,000 tokens:  $\sim 700$  unique words
- After 10,000 tokens: not 7,000, but maybe  $\sim 2,500\text{--}3,500$

Why it matters?

# Heaps' law

As you read more tokens in a corpus, you keep seeing new word **types**, but the **rate** of new words **slows down**.

## Example

- After 1,000 tokens:  $\sim 700$  unique words
- After 10,000 tokens: not 7,000, but maybe  $\sim 2,500$ – $3,500$

## Why it matters?

- Estimate how much data you need before vocabulary “stabilizes”

# Heaps' law

As you read more tokens in a corpus, you keep seeing new word **types**, but the **rate** of new words **slows down**.

## Example


- After 1,000 tokens:  $\sim 700$  unique words
- After 10,000 tokens: not 7,000, but maybe  $\sim 2,500$ – $3,500$

## Why it matters?

- Estimate how much data you need before vocabulary “stabilizes”
- Reminds us that growth is **sublinear**

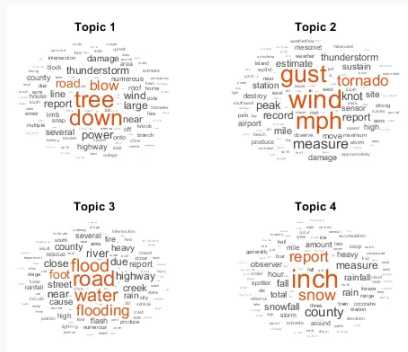
# Topic modeling

- A topic modeling is a type of **statistical modeling** for discovering the **abstract** topic that occur in a collection of documents.



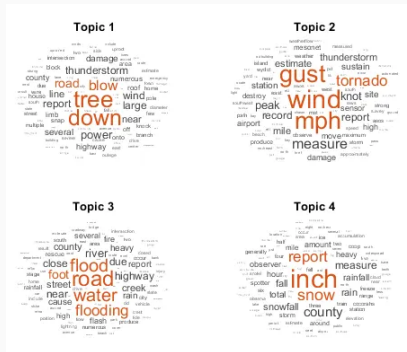
# Topic modeling

- A topic modeling is a type of **statistical modeling** for discovering the **abstract** topic that occur in a collection of documents.
- **Exploratory**: No annotated labels; discover latent structure using word frequencies/distributions



# Topic modeling

- A topic modeling is a type of **statistical modeling** for discovering the **abstract** topic that occur in a collection of documents.
- **Exploratory**: No annotated labels; discover latent structure using word frequencies/distributions
- Classic model: LDA (Blei et al., 2003).



Key idea: (1) Each document is a *mixture of topics*. (2) Each topic is a *distribution over words*. (3) Given only the words, LDA uses Bayesian inference to approximate the hidden topic structure.

1. Collect documents; tokenize, lemmatize; remove stop words.



Key idea: (1) Each document is a *mixture of topics*. (2) Each topic is a *distribution over words*. (3) Given only the words, LDA uses Bayesian inference to approximate the hidden topic structure.

1. Collect documents; tokenize, lemmatize; remove stop words.
2. Choose  $K$  topics; start with random assignments.

Key idea: (1) Each document is a *mixture of topics*. (2) Each topic is a *distribution over words*. (3) Given only the words, LDA uses Bayesian inference to approximate the hidden topic structure.

1. Collect documents; tokenize, lemmatize; remove stop words.
2. Choose  $K$  topics; start with random assignments.
3. Iterate: update topic assignments using document–topic and topic–word counts.

Key idea: (1) Each document is a *mixture of topics*. (2) Each topic is a *distribution over words*. (3) Given only the words, LDA uses Bayesian inference to approximate the hidden topic structure.

1. Collect documents; tokenize, lemmatize; remove stop words.
2. Choose  $K$  topics; start with random assignments.
3. Iterate: update topic assignments using document–topic and topic–word counts.
4. Inspect top words per topic; give each topic a human-readable label.

# LDA: minimal workflow

Key idea: (1) Each document is a *mixture of topics*. (2) Each topic is a *distribution over words*. (3) Given only the words, LDA uses Bayesian inference to approximate the hidden topic structure.

1. Collect documents; tokenize, lemmatize; remove stop words.
2. Choose  $K$  topics; start with random assignments.
3. Iterate: update topic assignments using document–topic and topic–word counts.
4. Inspect top words per topic; give each topic a human-readable label.
5. We'll do some hands-on practice with topic modeling.

## Word vectors

---

- Words themselves cannot be given as inputs to computers

- Words themselves cannot be given as inputs to computers
- BUT, numbers can be given as inputs to computers


- Words themselves cannot be given as inputs to computers
- BUT, numbers can be given as inputs to computers
- Encoding = converting words to vectors



- Words themselves cannot be given as inputs to computers
- BUT, numbers can be given as inputs to computers
- Encoding = converting words to vectors
  - *vector*: an ordered list of numbers (e.g., [0.1, 0.3, -0.5])


# Word vectors

- If two words often occur near the same neighbors (e.g., *dog* and *cat* near *cute*), their vectors will be similar.



# Word vectors

- If two words often occur near the same neighbors (e.g., *dog* and *cat* near *cute*), their vectors will be similar.
- *How?* Algorithms can automatically learn these vectors from corpus data



Core idea:

- Start with a large corpus

Core idea:

- Start with a large corpus
- Every word in a fixed vocabulary is represented by a vector

Core idea:

- Start with a large corpus
- Every word in a fixed vocabulary is represented by a vector
- Go through each position  $t$  in the text, which has a center word and a context word

Core idea:

- Start with a large corpus
- Every word in a fixed vocabulary is represented by a vector
- Go through each position  $t$  in the text, which has a center word and a context word
- Calculate the probability of a center word given a context word (or vice versa)

Core idea:

- Start with a large corpus
- Every word in a fixed vocabulary is represented by a vector
- Go through each position  $t$  in the text, which has a center word and a context word
- Calculate the probability of a center word given a context word (or vice versa)
- Keep adjusting the word vectors to **maximize** the probability



# Word vectors

Core idea:

- Start with a large corpus
- Every word in a fixed vocabulary is represented by a vector
- Go through each position  $t$  in the text, which has a center word and a context word
- Calculate the probability of a center word given a context word (or vice versa)
- Keep adjusting the word vectors to **maximize** the probability
- (*more on this in the NLP class!*)

- Once the vectors are learned, word vectors can be used for **mathematical operations**.

- Once the vectors are learned, word vectors can be used for **mathematical operations**.
- *Word2Vec* (Mikolov et al. 2013):

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

- Once the vectors are learned, word vectors can be used for **mathematical operations**.
- *Word2Vec* (Mikolov et al. 2013):

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

- We'll also explore *Word2Vec*.

# Submission/Grading Guidelines

- **Text/Docx file** (either `.txt` or `.docx`): Submit your output from the corpus exploration on Tuesday. (If you missed class, complete it individually and submit.) (10 points)
- **Notebook file** (`.ipynb`): Submit your work from today's session.
  - Topic Modeling (5 points)
    - + 5 extra points, if you experiment this code on another corpus
  - Word2Vec (5 points)
- **Optional:** Please complete the *Collocation* Tutorial for extra credit (+3 points above max)
  - Guidelines/Code are on the last week's section (course website)
- **PLEASE run all the codes, so the grader can seamlessly check your outputs!**