

15. Chatbot

LING-351 Language Technology and LLMs

Instructor: Hakyung Sung

October 21, 2025

Table of contents

1. What is a Chatbot?
2. How to build a chatbot?
3. Evaluation
4. Wrap-up

What is a Chatbot?

- A **dialog system** designed to engage in natural, sociable conversation with a human user.

- A **dialog system** designed to engage in natural, sociable conversation with a human user.
- Can support both **text-based** (written) and **speech-based** (spoken) interaction.

- A **dialog system** designed to engage in natural, sociable conversation with a human user.
- Can support both **text-based** (written) and **speech-based** (spoken) interaction.
- Typically involves modules for:

- A **dialog system** designed to engage in natural, sociable conversation with a human user.
- Can support both **text-based** (written) and **speech-based** (spoken) interaction.
- Typically involves modules for:
 - Speech-to-Text (STT)

- A **dialog system** designed to engage in natural, sociable conversation with a human user.
- Can support both **text-based** (written) and **speech-based** (spoken) interaction.
- Typically involves modules for:
 - Speech-to-Text (STT)
 - Text-to-Speech (TTS)

- The chatbot should respond *appropriately* to the user's input.

- The chatbot should respond *appropriately* to the user's input.
- **What does “appropriately” mean?**

- The chatbot should respond *appropriately* to the user's input.
- **What does “appropriately” mean?**
 - Contextually relevant

- The chatbot should respond *appropriately* to the user's input.
- **What does “appropriately” mean?**
 - Contextually relevant
 - Coherent and natural in tone

- The chatbot should respond *appropriately* to the user's input.
- **What does “appropriately” mean?**
 - Contextually relevant
 - Coherent and natural in tone
 - Aligned with conversational goals

- The chatbot should respond *appropriately* to the user's input.
- **What does “appropriately” mean?**
 - Contextually relevant
 - Coherent and natural in tone
 - Aligned with conversational goals
 - Anything else?

Why do we need a chatbot?

“a chatbot might serve as a digital companion”

Movie *Her* (2013)

<https://www.youtube.com/watch?v=f9Hg1x-Ctlw>

Group discussion

(*Shared deck*)

Think about your own experiences with chatbots. Some prompting questions:

- Have you used any specific platforms (e.g., ChatGPT, Grok, Replika)?
- When did the interaction feel *good* or *bad*?
- What made the chatbot seem more (or less) like a genuine companion?
- Was it the language style, emotional tone, or responsiveness?
- What do you think was the chatbot's main purpose in that interaction?
- How does this relate to the idea of a "digital companion" in the movie *Her*?
- If you haven't had any experience yet, why do you think that is?

- Whether or not a chatbot can truly provide fulfilling companionship, building one serves as an intellectual exploration of artificial intelligence (AI).

- Whether or not a chatbot can truly provide fulfilling companionship, building one serves as an intellectual exploration of artificial intelligence (AI).
- Language is often regarded as a key component in achieving artificial general intelligence (AGI)—the hypothetical ability of a machine to act with human-like *sentience* and *self-awareness*.

Turing test

- British mathematician *Alan Turing* (1950) proposed the Turing Test as a way to evaluate a machine's intelligence.

Turing test

- British mathematician *Alan Turing* (1950) proposed the Turing Test as a way to evaluate a machine's intelligence.
- The test asks whether a computer can convince a human interlocutor that it, too, is human during a text-based conversation.

Turing test

- British mathematician *Alan Turing* (1950) proposed the Turing Test as a way to evaluate a machine's intelligence.
- The test asks whether a computer can convince a human interlocutor that it, too, is human during a text-based conversation.
- If the human cannot reliably tell the difference, the machine is said to have demonstrated *human-like intelligence*.

Turing test

- British mathematician *Alan Turing* (1950) proposed the Turing Test as a way to evaluate a machine's intelligence.
- The test asks whether a computer can convince a human interlocutor that it, too, is human during a text-based conversation.
- If the human cannot reliably tell the difference, the machine is said to have demonstrated *human-like intelligence*.
- <https://www.youtube.com/watch?v=qLbYrvV7Ma4>

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?
- Several chatbots have “passed” the test:

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?
- Several chatbots have “passed” the test:
 - **Eugene Goostman** — pretended to be a 14-year-old Ukrainian learning English.

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?
- Several chatbots have “passed” the test:
 - **Eugene Goostman** — pretended to be a 14-year-old Ukrainian learning English.
 - **ELIZA** — mirrored the user’s words to make them feel understood.


Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?
- Several chatbots have “passed” the test:
 - **Eugene Goostman** — pretended to be a 14-year-old Ukrainian learning English.
 - **ELIZA** — mirrored the user’s words to make them feel understood.
- These cases reveal more about the limitations of the Turing Test than the intelligence of the chatbots.

Turin test - Controversy

- The Turing Test is controversial because humans themselves vary widely.
- What kind of human should a chatbot “pass” as—a child, a teenager, or an expert adult?
- Several chatbots have “passed” the test:
 - **Eugene Goostman** — pretended to be a 14-year-old Ukrainian learning English.
 - **ELIZA** — mirrored the user’s words to make them feel understood.
- These cases reveal more about the limitations of the Turing Test than the intelligence of the chatbots.
- <https://www.youtube.com/watch?v=3wLqsRLvV-c>

More thoughts



For any given chatbot, it's often unclear whether it is meant to be a social companion, a scientific experiment, or both—or whether the purpose is left for the **user to decide**.

How to build a chatbot?

Several ways to build a chatbot

1. Brute force
2. Rule-based
3. Corpus-trained
4. Language generation task

1. Brute-Force Chatbot

- A **brute-force chatbot** is built by **manually scripting responses** to a large set of expected user questions.

1. Brute-Force Chatbot

- A **brute-force chatbot** is built by **manually scripting responses** to a large set of expected user questions.
- Developers brainstorm the most common queries (e.g., “When were you born?”) and write pre-defined answers that reflect the chatbot’s personality or brand voice.

1. Brute-Force Chatbot

- A **brute-force chatbot** is built by **manually scripting responses** to a large set of expected user questions.
- Developers brainstorm the most common queries (e.g., “When were you born?”) and write pre-defined answers that reflect the chatbot’s personality or brand voice.
- This ensures reasonable replies for predictable inputs—like Siri knowing its launch year— but fails to handle unexpected or creative questions.

1. Brute-Force Chatbot

- A **brute-force chatbot** is built by **manually scripting responses** to a large set of expected user questions.
- Developers brainstorm the most common queries (e.g., “When were you born?”) and write pre-defined answers that reflect the chatbot’s personality or brand voice.
- This ensures reasonable replies for predictable inputs—like Siri knowing its launch year— but fails to handle unexpected or creative questions.
- It mimics understanding through memorization rather than reasoning.

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like
 - If the input contains “my X” → reply “Tell me more about your X.”

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like
 - If the input contains "my X" → reply "Tell me more about your X."
 - If the input contains "I am X" → reply "How long have you been X?"

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like
 - If the input contains "my X" → reply "Tell me more about your X."
 - If the input contains "I am X" → reply "How long have you been X?"
 - <https://www.youtube.com/watch?v=RMK9AphfLco>

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like
 - If the input contains "my X" → reply "Tell me more about your X."
 - If the input contains "I am X" → reply "How long have you been X?"
 - <https://www.youtube.com/watch?v=RMK9AphfLco>
- By asking questions rather than making statements, ELIZA maintains surface-level relevance without true understanding.

2. Rule-Based Chatbot

- A **rule-based chatbot** generates responses using **predefined if-then rules** written by humans.
- Example: **ELIZA** (Weizenbaum, 1966) used patterns like
 - If the input contains "my X" → reply "Tell me more about your X."
 - If the input contains "I am X" → reply "How long have you been X?"
 - <https://www.youtube.com/watch?v=RMK9AphfLco>
- By asking questions rather than making statements, ELIZA maintains surface-level relevance without true understanding.
- Like brute-force systems, rule-based chatbots lack awareness of self, user, or world, and their rule sets quickly become complex, rigid, and incomplete.

3. Corpus-Trained Chatbot

- A **corpus-trained chatbot** is built by **training on large collections of real dialog and text** (e.g., film scripts, TV subtitles, Wikipedia).

3. Corpus-Trained Chatbot

- A **corpus-trained chatbot** is built by **training on large collections of real dialog and text** (e.g., film scripts, TV subtitles, Wikipedia).
- It treats conversation as a search problem — selecting the one or two lines from the corpus most similar to the user's preceding utterance.

3. Corpus-Trained Chatbot

- A **corpus-trained chatbot** is built by **training on large collections of real dialog and text** (e.g., film scripts, TV subtitles, Wikipedia).
- It treats conversation as a search problem — selecting the one or two lines from the corpus most similar to the user's preceding utterance.
- This approach aims for conversational relevance and maintains quality, quantity, and manner by reusing well-formed, contextually appropriate sentences.

3. Corpus-Trained Chatbot

- A **corpus-trained chatbot** is built by **training on large collections of real dialog and text** (e.g., film scripts, TV subtitles, Wikipedia).
- It treats conversation as a search problem — selecting the one or two lines from the corpus most similar to the user's preceding utterance.
- This approach aims for conversational relevance and maintains quality, quantity, and manner by reusing well-formed, contextually appropriate sentences.
- The chatbot's "knowledge" of itself, the user, and the world is limited to what is encoded in the corpus.

4. Generative Chatbot

- A **generative chatbot** treats conversation as a **language generation task**, using machine learning models to predict and produce the next most likely utterance.

4. Generative Chatbot

- A **generative chatbot** treats conversation as a **language generation task**, using machine learning models to predict and produce the next most likely utterance.
- Unlike corpus-trained systems that retrieve lines, it creates new responses by distilling patterns across massive text corpora.

4. Generative Chatbot

- A **generative chatbot** treats conversation as a **language generation task**, using machine learning models to predict and produce the next most likely utterance.
- Unlike corpus-trained systems that retrieve lines, it creates new responses by distilling patterns across massive text corpora.
- Modern generative chatbots still rely partly on **brute-force scripts** for sensitive or restricted topics, and may use a scripted “autobiography” to define their identity.

4. Generative Chatbot

- A **generative chatbot** treats conversation as a **language generation task**, using machine learning models to predict and produce the next most likely utterance.
- Unlike corpus-trained systems that retrieve lines, it creates new responses by distilling patterns across massive text corpora.
- Modern generative chatbots still rely partly on **brute-force scripts** for sensitive or restricted topics, and may use a scripted “autobiography” to define their identity.
- They represent the world, the user, and prior context through the information encoded in their training data.

Evaluation

- **Evaluation begins with purpose:**

- **Evaluation begins with purpose:**
 - A chatbot for companionship is judged by user engagement or satisfaction

- **Evaluation begins with purpose:**

- A chatbot for companionship is judged by user engagement or satisfaction
- for intelligence is tested by human-likeness or reasoning ability.

- **Types of evaluation:**

- **Types of evaluation:**
 - **Informal:** exploratory, subjective — like an interview.

- **Types of evaluation:**

- **Informal:** exploratory, subjective — like an interview.
- **Formal:** systematic, objective — like an exam, with fixed questions and scoring criteria.

Evaluating Chatbots

- **Dimensions to test:**

Evaluating Chatbots

- **Dimensions to test:**
 - **Self-concept:** consistency across multiple turns.

Evaluating Chatbots

- **Dimensions to test:**
 - **Self-concept:** consistency across multiple turns.
 - **Common Ground:** memory of prior context.

- **Dimensions to test:**
 - **Self-concept:** consistency across multiple turns.
 - **Common Ground:** memory of prior context.
 - **Interlocutor model:** understanding implied meaning.

Evaluating Chatbots

- **Dimensions to test:**
 - **Self-concept:** consistency across multiple turns.
 - **Common Ground:** memory of prior context.
 - **Interlocutor model:** understanding implied meaning.
 - **World knowledge:** reasoning about everyday facts (e.g., "Can a tennis ball fit into a toaster?").

Wrap-up

Wrap-up

- What is a chatbot?

Wrap-up

- What is a chatbot?
- Different ways to build them

Wrap-up

- What is a chatbot?
- Different ways to build them
- Evaluation

Wrap-up

- What is a chatbot?
- Different ways to build them
- Evaluation
- On Thursday: We will build a simple chatbot using RAG (Retrieval-Augmented Generation).